

# Diabetes Prediction using classification method

## Import Libraries

```
In [84]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
```

## Load Datasets

```
In [85]: data = pd.read_csv("diabetes.csv")
```

## Shape of Data

```
In [86]: data.shape
```

```
Out[86]: (768, 9)
```

## First 8 Rows of Data

```
In [87]: data.head(8)
```

```
Out[87]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction
0	6	148	72	35	0	33.6	0.627
1	1	85	66	29	0	26.6	0.351
2	8	183	64	0	0	23.3	0.672
3	1	89	66	23	94	28.1	0.167
4	0	137	40	35	168	43.1	2.286
5	5	116	74	0	0	25.6	0.201
6	3	78	50	32	88	31.0	0.246
7	10	115	0	0	0	35.3	0.134

## Last 7 Rows of Data

In [88]: `data.tail(7)`

Out[88]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunc
<b>761</b>	9	170	74	31	0	44.0	0.4
<b>762</b>	9	89	62	0	0	22.5	0.1
<b>763</b>	10	101	76	48	180	32.9	0.1
<b>764</b>	2	122	70	27	0	36.8	0.3
<b>765</b>	5	121	72	23	112	26.2	0.2
<b>766</b>	1	126	60	0	0	30.1	0.3
<b>767</b>	1	93	70	31	0	30.4	0.3

In [89]: `data.describe()`

Out[89]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabete
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	
<b>mean</b>	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	
<b>std</b>	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	
<b>min</b>	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	
<b>50%</b>	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	

In [90]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [91]: data.value\_counts()

```
Out[91]: Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  DiabetesP
edegreeFunction  Age  Outcome
0      57      60      0      0      21.7  0.735
67  0      1
67      67      76      0      0      45.3  0.194
46  0      1
5      103      108      37      0      39.2  0.305
65  0      1
65      104      74      0      0      28.8  0.153
48  0      1
28  0      105      72      29      325      36.9  0.159
28  0      1
..
2      84      50      23      76      30.4  0.968
21  0      1
21      85      65      0      0      39.6  0.930
27  0      1
27      87      0      23      0      28.9  0.773
25  0      1
25      58      16      52      32.7  0.166
25  0      1
17      163      72      41      114      40.9  0.817
47  1      1
Name: count, Length: 768, dtype: int64
```

In [92]: data.columns

```
Out[92]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
               'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
              dtype='object')
```

## Checking Null Values


```
In [93]: data.isnull().sum()
```

```
Out[93]: Pregnancies      0
          Glucose          0
          BloodPressure    0
          SkinThickness     0
          Insulin           0
          BMI               0
          DiabetesPedigreeFunction  0
          Age              0
          Outcome           0
          dtype: int64
```

## Diabetes Distribution

```
In [94]: #Finding Class Distribution Percentage
print(data['Outcome'].value_counts(ascending=True))
print(data['Outcome'].value_counts(1,ascending=True).apply(lambda x: format(x
print()
# Plot the bar chart
data['Outcome'].value_counts(normalize=True).plot(kind='barh',figsize=(10, 2))
plt.title('Diabetes Distribution (%)', fontsize=18)
plt.yticks(ticks=[0,1], labels=['Non-Diabetic', 'Diabetic'])
plt.show()
```

```
Outcome
1    268
0    500
Name: count, dtype: int64
Outcome
1    34.895833%
0    65.104167%
Name: proportion, dtype: object
```



Exploratory Data Analysis

In [95]:

data.corr()


Out[95]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin		
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.01	
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.22	
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.28	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.39	
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.19	
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.00	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	0.185071	0.14	
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.03	
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.29	

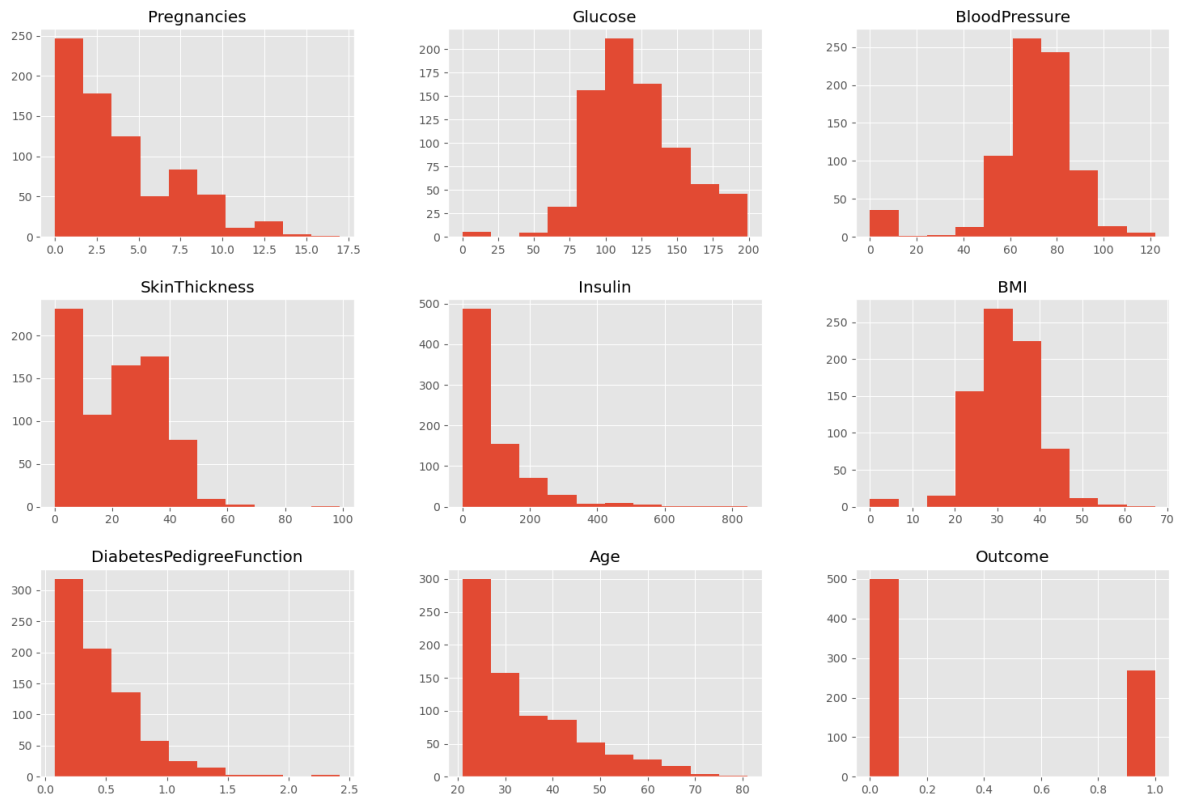
## Correlation Matrix

```
In [96]: plt.figure(figsize = (12,10))  
  
sns.heatmap(data.corr(), annot = True)
```

Out[96]: <Axes: >




```
In [97]: data.hist(figsize=(18,12))  
plt.show()
```



```
In [98]: features = ['Glucose', 'BloodPressure', 'Insulin', 'BMI', 'Age', 'SkinThickness']
plt.figure(figsize=(14, 10))

for i, feature in enumerate(features, start=1):
    plt.subplot(2, 3, i)
    sns.boxplot(x=feature, data=data)

plt.tight_layout()
plt.show()
```






```
In [99]: import seaborn as sns
import matplotlib.pyplot as plt
mean_col = ['Glucose', 'BloodPressure', 'Insulin', 'Age', 'Outcome', 'BMI']
sns.pairplot(data[mean_col])
plt.show()
```

C:\Users\Dell\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)




```
In [100]: sns.boxplot(x='Outcome',y='Insulin',data=data)
```

```
Out[100]: <Axes: xlabel='Outcome', ylabel='Insulin'>
```



```
In [101]: sns.regplot(x='BMI', y='Glucose', data=data)
```


```
Out[101]: <Axes: xlabel='BMI', ylabel='Glucose'>
```



```
In [102]: sns.relplot(x='BMI', y= 'Glucose', data=data)
```


C:\Users\Dell\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)

```
Out[102]: <seaborn.axisgrid.FacetGrid at 0x1b7c8ad2050>
```




```
In [103]: sns.scatterplot(x='Glucose', y= 'Insulin', data=data)
```

```
Out[103]: <Axes: xlabel='Glucose', ylabel='Insulin'>
```



```
In [104]: sns.jointplot(x='SkinThickness', y= 'Insulin', data=data)
```

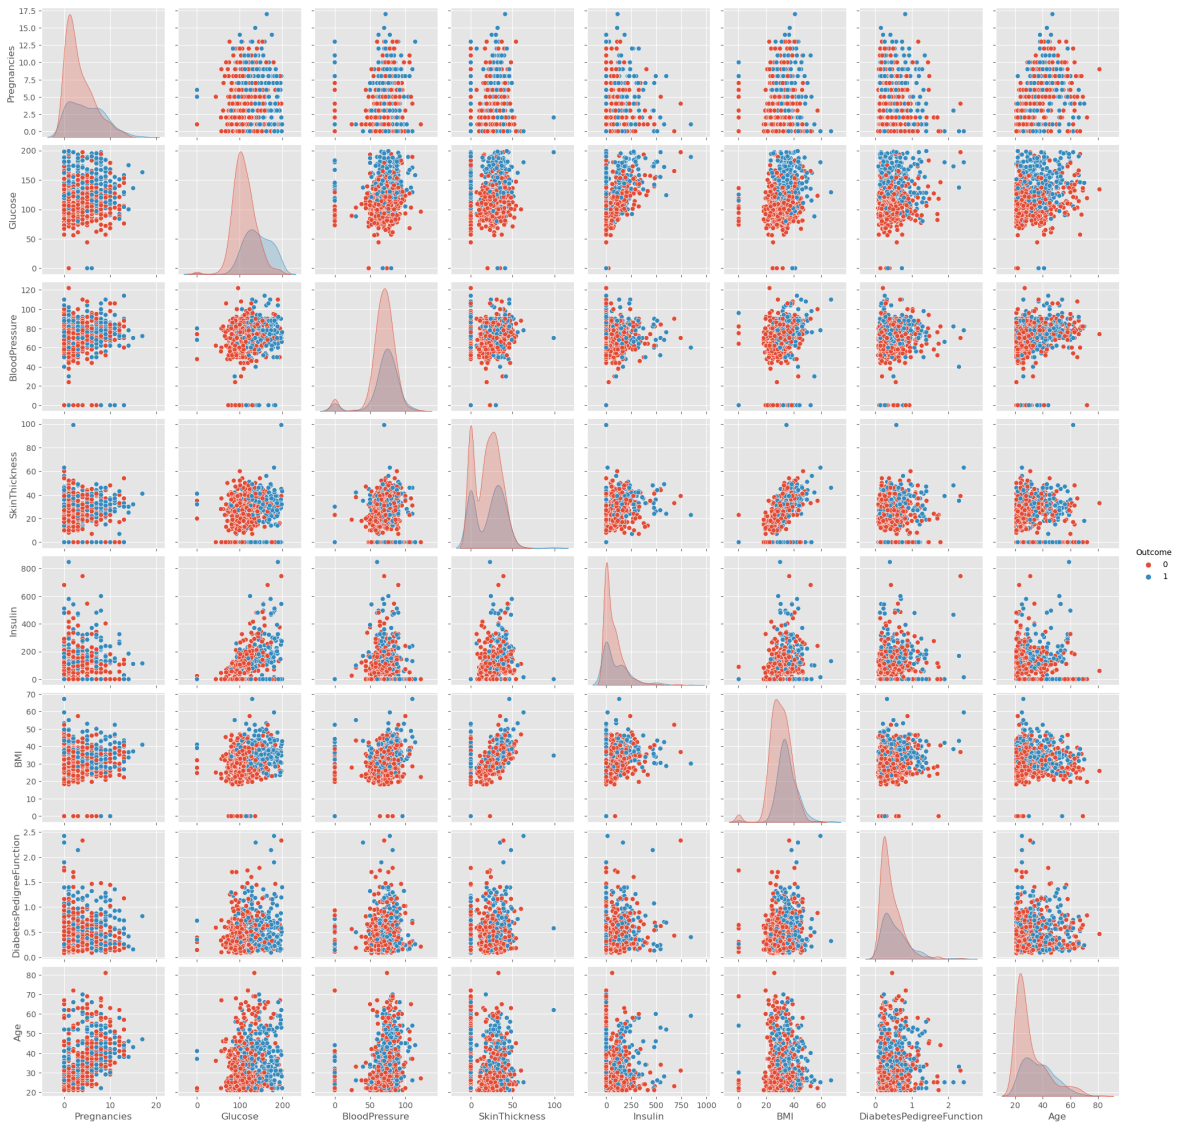
```
Out[104]: <seaborn.axisgrid.JointGrid at 0x1b7ca5450d0>
```



```
In [105]: sns.pairplot(data,hue='Outcome')
```


```
C:\Users\Dell\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
```

```
Out[105]: <seaborn.axisgrid.PairGrid at 0x1b7c6f9bc10>
```



```
In [106]: sns.lineplot(x='Glucose', y='Insulin', data=data)
```


```
Out[106]: <Axes: xlabel='Glucose', ylabel='Insulin'>
```






```
In [107]: sns.stripplot(x='Glucose', y='Insulin', data=data, jitter=True, size=3)
```


```
Out[107]: <Axes: xlabel='Glucose', ylabel='Insulin'>
```



```
In [108]: sns.barplot(x="SkinThickness", y="Insulin", data=data[150:180])  
plt.title("SkinThickness vs Insulin",fontsize=15)  
plt.xlabel("SkinThickness")  
plt.ylabel("Insulin")  
plt.show()  
plt.style.use("ggplot")
```



```
In [109]: plt.figure(figsize=(5,5))
sns.barplot(x="Glucose", y="Insulin", data=data[120:130])
plt.title("Glucose vs Insulin",fontsize=15)
plt.xlabel("Glucose")
plt.ylabel("Insulin")
plt.show()
```



## ## Pre-process, Training and Testing Data

```
In [123]: x = data.drop(columns = 'Outcome')

y = data['Outcome']

from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test = train_test_split(x,y,test_size=0.3,random_star
```

## Train the Neural Network model

```
In [124]: from sklearn.neural_network import MLPClassifier
nn_model = MLPClassifier(hidden_layer_sizes=(100,50), max_iter=500, activation
nn_model.fit(X_train, y_train)
```

```
Out[124]:
MLPClassifier
MLPClassifier(hidden_layer_sizes=(100, 50), max_iter=500, random_state=42)
```

## MODELS

### 1. Logistic Regression

```
In [125]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

from sklearn.metrics import accuracy_score
LRAcc = accuracy_score(y_pred,y_test)
print('Logistic Regression accuracy is: {:.2f}%'.format(LRAcc*100))
```

	precision	recall	f1-score	support
0	0.80	0.80	0.80	151
1	0.62	0.62	0.62	80
accuracy			0.74	231
macro avg	0.71	0.71	0.71	231
weighted avg	0.74	0.74	0.74	231

```
[[121  30]
 [ 30  50]]
```

Logistic Regression accuracy is: 74.03%

C:\Users\Dell\anaconda3\Lib\site-packages\sklearn\linear\_model\\_logistic.py:  
460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression) ([https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression))

```
n_iter_i = _check_optimize_result(
```

## 2. SVM

```
In [126]: from sklearn.svm import SVC # Correct import statement
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

model = SVC()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

SVMacc = accuracy_score(y_test, y_pred)
print('SVM accuracy is: {:.2f}%'.format(SVMacc * 100))
```

	precision	recall	f1-score	support
0	0.76	0.87	0.81	151
1	0.66	0.49	0.56	80
accuracy			0.74	231
macro avg	0.71	0.68	0.69	231
weighted avg	0.73	0.74	0.72	231

```
[[131  20]
 [ 41  39]]
SVM accuracy is: 73.59%
```

## 3. Decision Tree

```
In [127]: from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score

model = DecisionTreeClassifier()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))

DTAcc = accuracy_score(y_test, y_pred)
print('Decision Tree accuracy is: {:.2f}%'.format(DTAcc * 100))
```

	precision	recall	f1-score	support
0	0.79	0.71	0.75	151
1	0.54	0.64	0.58	80
accuracy			0.68	231
macro avg	0.66	0.67	0.66	231
weighted avg	0.70	0.68	0.69	231

[[107 44]  
[ 29 51]]  
Decision Tree accuracy is: 68.40%

## Compare Models

```
In [128]: compare = pd.DataFrame({'Models Trained': ['Logistic Regression', 'SVM', 'Decision Tree'],
                                'Accuracy': [LRAcc*100, SVMAcc*100, DTAcc*100]})
compare.sort_values(by='Accuracy', ascending=False)
```

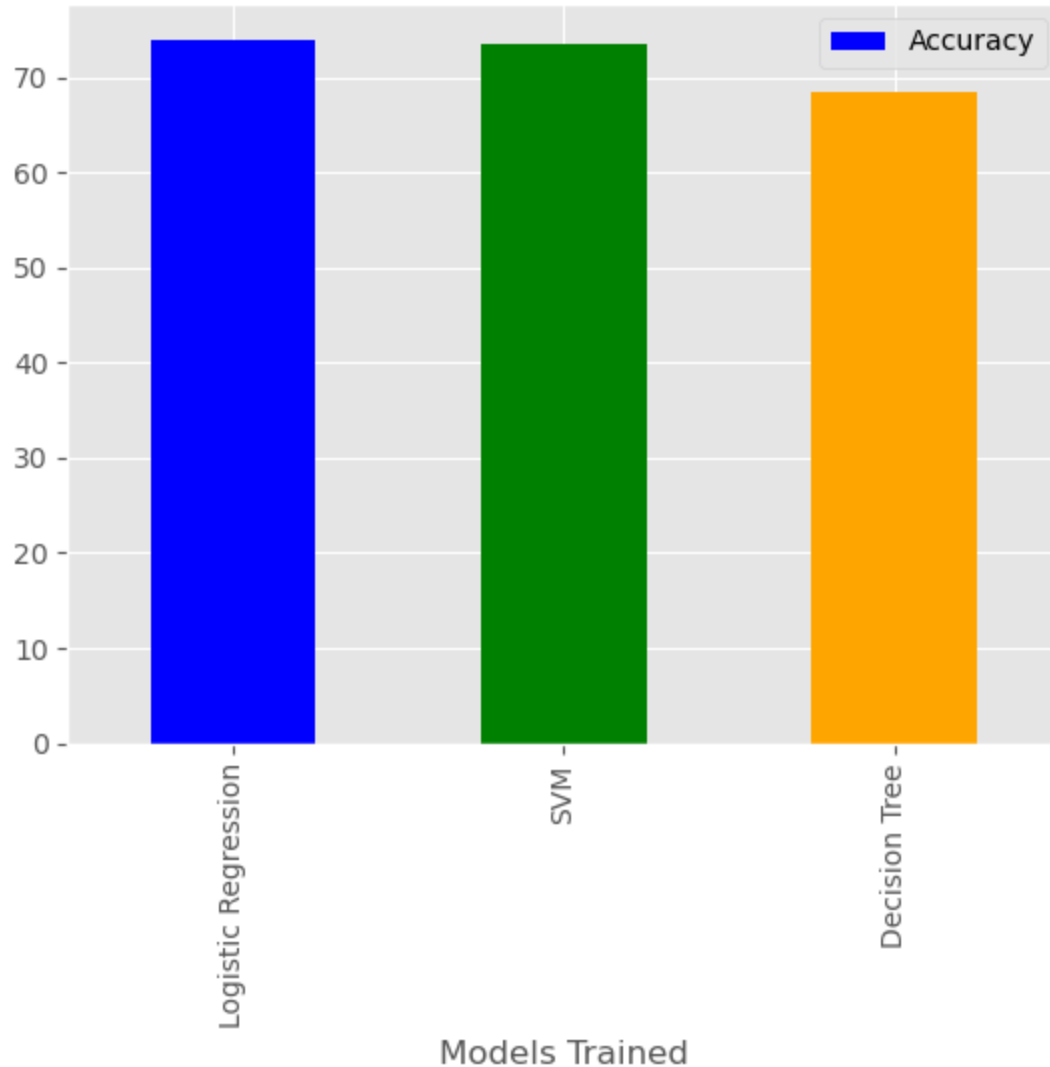
Out[128]:

	Models Trained	Accuracy
0	Logistic Regression	74.025974
1	SVM	73.593074
2	Decision Tree	68.398268

## Plotting Model Comparison

```
In [130]: compare.plot(x='Models Trained', y='Accuracy', kind='bar', color=['blue', 'green', 'orange'])
```

```
Out[130]: <Axes: xlabel='Models Trained'>
```



**From the comparison plot, among the 3 Machine Learning Models, Logistic Regression had achieved the highest accuracy of 74.025974%.**